

# DISTRIBUTED SYSTEMS

## *Concepts and Design*

Fifth Edition



George Coulouris  
Jean Dollimore  
Tim Kindberg  
Gordon Blair



PEARSON

# DISTRIBUTED SYSTEMS

## *Concepts and Design*

Fifth Edition

*This page intentionally left blank*

# DISTRIBUTED SYSTEMS

## *Concepts and Design*

Fifth Edition

George Coulouris  
*Cambridge University*

Jean Dollimore  
*formerly of Queen Mary,  
University of London*

Tim Kindberg  
*matter 2 media*

Gordon Blair  
*Lancaster University*

International Edition  
contributions by:  
Arup Kumar  
Bhattacharjee  
*RCC Institute of Information  
Technology, Kolkata*

Soumen Mukherjee  
*RCC Institute of Information  
Technology, Kolkata*

**Addison-Wesley**

Boston Columbus Indianapolis New York San Francisco Upper Saddle River  
Amsterdam Cape Town Dubai London Madrid Milan Munich Paris Montréal Toronto  
Delhi Mexico City São Paulo Sydney Hong Kong Seoul Singapore Taipei Tokyo

Editorial Director: Marcia Horton  
Editor-in-Chief: Michael Hirsch  
Executive Editor: Matt Goldstein  
Editorial Assistant: Chelsea Bell  
Vice President, Marketing: Patrice Jones  
Marketing Manager: Yezan Alayan  
Marketing Coordinator: Kathryn Ferranti  
Vice President, Production: Vince O'Brien  
Managing Editor: Jeff Holcomb  
Senior Production Project Manager: Marilyn Lloyd  
Publisher, International Edition: Angshuman Chakraborty  
Acquisitions Editor, International Edition: Arunabha Deb  
Publishing Assistant, International Edition: Shokhi Shah  
Senior Operations Supervisor: Alan Fischer  
Manufacturing Buyer: Lisa McDowell  
Art Director: Jayne Conte  
Cover Designer: Suzanne Duda  
Cover Image: Sky: © amygdala\_imagery; Kite: © Alamy; Mobile phone: © yasinguneyu/iStock  
Media Editor: Daniel Sandin  
Media Project Manager: Wanda Rockwell  
Cover Printer: Lehigh-Phoenix Color

---

Pearson Education Limited  
Edinburgh Gate  
Harlow  
Essex CM20 2JE  
England

and Associated Companies throughout the world

Visit us on the World Wide Web at:  
[www.pearsoninternationaleditions.com](http://www.pearsoninternationaleditions.com)

© Pearson Education Limited 2012

The rights of George Coulouris, Jean Dollimore, Tim Kindberg and Gordon Blair to be identified as authors of this work have been asserted by them in accordance with the Copyright, Designs and Patents Act 1988.

*Authorized adaptation from the United States edition, entitled Distributed Systems, Concepts and Design, 5<sup>th</sup> edition, ISBN 978-0-13-214301-1 by George Coulouris, Jean Dollimore, Tim Kindberg and Gordon Blair published by Pearson Education © 2012.*

All rights reserved. No part of this publication may be reproduced, stored in a retrieval system, or transmitted in any form or by any means, electronic, mechanical, photocopying, recording or otherwise, without either the prior written permission of the publisher or a licence permitting restricted copying in the United Kingdom issued by the Copyright Licensing Agency Ltd, Saffron House, 6–10 Kirby Street, London EC1N 8TS.

All trademarks used herein are the property of their respective owners. The use of any trademark in this text does not vest in the author or publisher any trademark ownership rights in such trademarks, nor does the use of such trademarks imply any affiliation with or endorsement of this book by such owners.

Microsoft® and Windows® are registered trademarks of the Microsoft Corporation in the U.S.A. and other countries. Screen shots and icons reprinted with permission from the Microsoft Corporation. This book is not sponsored or endorsed by or affiliated with the Microsoft Corporation.

British Library Cataloguing-in-Publication Data

A catalogue record for this book is available from the British Library

10 9 8 7 6 5 4 3 2 1  
14 13 12 11 10

Typeset in Times New Roman by authors using FrameMaker

Printed and bound by Edwards Brothers in The United States of America

The publisher's policy is to use paper manufactured from sustainable forests.

ISBN 10: 0-273-76059-9

ISBN 13: 978-0-273-76059-7

**Addison-Wesley**  
is an imprint of



[www.pearsonhighered.com](http://www.pearsonhighered.com)

<b>4</b>	<b>INTERPROCESS COMMUNICATION</b>	<b>161</b>
4.1	Introduction	162
4.2	The API for the Internet protocols	163
4.3	External data representation and marshalling	174
4.4	Multicast communication	185
4.5	Network virtualization: Overlay networks	190
4.6	Case study: MPI	194
4.7	Summary	197
<b>5</b>	<b>REMOTE INVOCATION</b>	<b>201</b>
5.1	Introduction	202
5.2	Request-reply protocols	203
5.3	Remote procedure call	211
5.4	Remote method invocation	220
5.5	Case study: Java RMI	233
5.6	Summary	241
<b>6</b>	<b>INDIRECT COMMUNICATION</b>	<b>245</b>
6.1	Introduction	246
6.2	Group communication	248
6.3	Publish-subscribe systems	258
6.4	Message queues	270
6.5	Shared memory approaches	278
6.6	Summary	290
<b>7</b>	<b>OPERATING SYSTEM SUPPORT</b>	<b>295</b>
7.1	Introduction	296
7.2	The operating system layer	297
7.3	Protection	300
7.4	Processes and threads	302
7.5	Communication and invocation	319
7.6	Operating system architecture	330
7.7	Virtualization at the operating system level	334
7.8	Summary	347

---

<b>8</b>	<b>DISTRIBUTED OBJECTS AND COMPONENTS</b>	<b>351</b>
8.1	Introduction	352
8.2	Distributed objects	353
8.3	Case study: CORBA	356
8.4	From objects to components	374
8.5	Case studies: Enterprise JavaBeans and Fractal	380
8.6	Summary	394
<b>9</b>	<b>WEB SERVICES</b>	<b>397</b>
9.1	Introduction	398
9.2	Web services	400
9.3	Service descriptions and IDL for web services	416
9.4	A directory service for use with web services	420
9.5	XML security	422
9.6	Coordination of web services	427
9.7	Applications of web services	429
9.8	Summary	435
<b>10</b>	<b>PEER-TO-PEER SYSTEMS</b>	<b>439</b>
10.1	Introduction	440
10.2	Napster and its legacy	444
10.3	Peer-to-peer middleware	446
10.4	Routing overlays	449
10.5	Overlay case studies: Pastry, Tapestry	452
10.6	Application case studies: Squirrel, OceanStore, Ivy	465
10.7	Summary	474
<b>11</b>	<b>SECURITY</b>	<b>479</b>
11.1	Introduction	480
11.2	Overview of security techniques	488
11.3	Cryptographic algorithms	500
11.4	Digital signatures	509
11.5	Cryptography pragmatics	516
11.6	Case studies: Needham–Schroeder, Kerberos, TLS, 802.11 WiFi	519
11.7	Summary	534

<b>12</b>	<b>DISTRIBUTED FILE SYSTEMS</b>	<b>537</b>
12.1	Introduction	538
12.2	File service architecture	546
12.3	Case study: Sun Network File System	552
12.4	Case study: The Andrew File System	564
12.5	Enhancements and further developments	573
12.6	Summary	579
<b>13</b>	<b>NAME SERVICES</b>	<b>581</b>
13.1	Introduction	582
13.2	Name services and the Domain Name System	585
13.3	Directory services	600
13.4	Case study: The Global Name Service	601
13.5	Case study: The X.500 Directory Service	604
13.6	Summary	608
<b>14</b>	<b>TIME AND GLOBAL STATES</b>	<b>611</b>
14.1	Introduction	612
14.2	Clocks, events and process states	613
14.3	Synchronizing physical clocks	615
14.4	Logical time and logical clocks	623
14.5	Global states	626
14.6	Distributed debugging	635
14.7	Summary	642
<b>15</b>	<b>COORDINATION AND AGREEMENT</b>	<b>645</b>
15.1	Introduction	646
15.2	Distributed mutual exclusion	649
15.3	Elections	657
15.4	Coordination and agreement in group communication	662
15.5	Consensus and related problems	675
15.6	Summary	687



---

<b>16</b>	<b>TRANSACTIONS AND CONCURRENCY CONTROL</b>	<b>691</b>
16.1	Introduction	692
16.2	Transactions	695
16.3	Nested transactions	706
16.4	Locks	708
16.5	Optimistic concurrency control	723
16.6	Timestamp ordering	727
16.7	Comparison of methods for concurrency control	734
16.8	Summary	736
<b>17</b>	<b>DISTRIBUTED TRANSACTIONS</b>	<b>743</b>
17.1	Introduction	744
17.2	Flat and nested distributed transactions	744
17.3	Atomic commit protocols	747
17.4	Concurrency control in distributed transactions	756
17.5	Distributed deadlocks	759
17.6	Transaction recovery	767
17.7	Summary	777
<b>18</b>	<b>REPLICATION</b>	<b>781</b>
18.1	Introduction	782
18.2	System model and the role of group communication	784
18.3	Fault-tolerant services	791
18.4	Case studies of highly available services: The gossip architecture, Bayou and Coda	798
18.5	Transactions with replicated data	818
18.6	Summary	830
<b>19</b>	<b>MOBILE AND UBIQUITOUS COMPUTING</b>	<b>833</b>
19.1	Introduction	834
19.2	Association	843
19.3	Interoperation	851
19.4	Sensing and context awareness	860
19.5	Security and privacy	873
19.6	Adaptation	882
19.7	Case study: Cooltown	887
19.8	Summary	894

---

<b>20</b>	<b>DISTRIBUTED MULTIMEDIA SYSTEMS</b>	<b>897</b>
20.1	Introduction	898
20.2	Characteristics of multimedia data	902
20.3	Quality of service management	903
20.4	Resource management	913
20.5	Stream adaptation	915
20.6	Case studies: Tiger, BitTorrent and End System Multicast	917
20.7	Summary	929
<b>21</b>	<b>DESIGNING DISTRIBUTED SYSTEMS: GOOGLE CASE STUDY</b>	<b>931</b>
21.1	Introduction	932
21.2	Introducing the case study: Google	933
21.3	Overall architecture and design philosophy	938
21.4	Underlying communication paradigms	944
21.5	Data storage and coordination services	951
21.6	Distributed computation services	972
21.7	Summary	980
	<b>REFERENCES</b>	<b>983</b>
	<b>INDEX</b>	<b>1041</b>